

Design and analysis of high performance conjugate gradient and Lanczos type algorithms

Tyler Chen

February XX, 2020

slides



`chen.pw/d/f5b`

code



`chen.pw/d/i7g`

proposal



`chen.pw/d/2mn`

Acknowledgements

We are gathered on the **unceded** land of the Coast Salish people, and in particular of the Duwamish Tribe.

Acknowledgements

This material is based upon work supported by the National Science Foundation Graduate Research Fellowship Program under Grant No. DGE-1762114. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author and do not necessarily reflect the views of the National Science Foundation.

Introduction

- Conjugate gradient (CG) is used to solve a linear system $\mathbf{Ax} = \mathbf{b}$ when $\mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric positive definite
- CG has low storage and floating point operation costs
 - one matrix vector product, two inner products, and a few vector updates each iteration
- Modern supercomputers have reached **exascale** (10^{18} flops)
 - Krylov subspace methods can only reach a fraction of this rate because of **communication costs**
- Convergence of CG in finite precision is not very well understood in general
- **Need to address communication costs, while considering numerical properties!**

The conjugate gradient algorithm

- At each step, CG generates iterate \mathbf{x}_k which minimizes the \mathbf{A} -norm of the error over,

$$\mathcal{K}_k(\mathbf{A}, \mathbf{r}_0) = \text{span}\{\mathbf{r}_0, \mathbf{A}\mathbf{r}_0, \dots, \mathbf{A}^{k-1}\mathbf{r}_0\}$$

- Builds up an \mathbf{A} -orthogonal basis of “search directions” $\mathbf{p}_0, \mathbf{p}_1, \mathbf{p}_2, \dots$ and minimizes independently along each

$$\langle \mathbf{e}_k, \mathbf{p}_j \rangle_{\mathbf{A}} = \langle \mathbf{A}^{-1}\mathbf{b} - \mathbf{x}_k, \mathbf{p}_j \rangle_{\mathbf{A}} = \langle \mathbf{b} - \mathbf{A}\mathbf{x}_k, \mathbf{p}_j \rangle = \langle \mathbf{r}_k, \mathbf{p}_k \rangle$$

- don't need to store all these search vectors because of **three term Lanczos recurrence!**

The conjugate gradient algorithm

Algorithm 1 Hestenes and Stiefel Conjugate Gradient (preconditioned)

```
1: procedure HS-CG( $\mathbf{A}$ ,  $\mathbf{M}$ ,  $\mathbf{b}$ ,  $\mathbf{x}_0$ )
2:   initialize()
3:   for  $k = 1, 2, \dots$  do
4:      $\mathbf{x}_k = \mathbf{x}_{k-1} + \alpha_{k-1}\mathbf{p}_{k-1}$ 
5:      $\mathbf{r}_k = \mathbf{r}_{k-1} - \alpha_{k-1}\mathbf{s}_{k-1}$ ,  $\tilde{\mathbf{r}}_k = \mathbf{M}^{-1}\mathbf{r}_k$ 
6:      $\nu_k = \langle \tilde{\mathbf{r}}_k, \mathbf{r}_k \rangle$ 
7:      $\beta_k = \nu_k / \nu_{k-1}$ 
8:      $\mathbf{p}_k = \tilde{\mathbf{r}}_k + \beta_k\mathbf{p}_{k-1}$ 
9:      $\mathbf{s}_k = \mathbf{A}\mathbf{p}_k$ 
10:     $\mu_k = \langle \mathbf{p}_k, \mathbf{s}_k \rangle$ 
11:     $\alpha_k = \nu_k / \mu_k$ 
12:  end for
13: end procedure
```

Lanczos

- Lanczos algorithm used to compute orthogonal basis $\mathbf{q}_0, \mathbf{q}_1, \dots$ for Krylov subspace $\mathcal{K}_k(\mathbf{A}, \mathbf{q}_0)$ when \mathbf{A} is symmetric
- produces three term recurrence

$$\mathbf{A}\mathbf{q}_j = \gamma_{j-1}\mathbf{q}_{j-1} + \delta_j\mathbf{q}_j + \gamma_j\mathbf{q}_{j+1}.$$

We can write this recurrence in matrix form as

$$\mathbf{A}\mathbf{Q}_k = \mathbf{Q}_k\mathbf{T}_k + \gamma_k\mathbf{q}_{k+1}\mathbf{e}_k^\top.$$

CG and Lanczos

- CG and Lanczos are equivalent; i.e. can be derived from one another.
- the columns of \mathbf{Q}_k are $\mathbf{q}_{j+1} = (-1)^j \frac{\mathbf{r}_j}{\|\mathbf{r}_j\|}$ and the diagonal and the off-diagonal entries of \mathbf{T}_k are respectively given by

$$\delta_j = \left(\frac{1}{\alpha_{j-1}} - \frac{\beta_{j-1}}{\alpha_{j-2}} \right), \quad \gamma_j = \frac{1}{\alpha_{j-1}} \frac{\|\mathbf{r}_j\|}{\|\mathbf{r}_{j-1}\|}.$$

Communication costs

- on large machines the cost of reading and moving data dominates the cost of floating point operations
- inner products and dense matrix vector products require global communication
- sparse matrix vector products require local communication
- vector updates require no communication

Hiding communication

- In each iteration we would like to **hide communication** by computing all inner products and matrix vector products simultaneously
- Do this by finding **mathematically equivalent** expressions for an inner product using recurrences
 - the new expressions are **not equivalent in finite precision**; i.e. the order of things has changed
- various ways of doing this have been studied before; see for instance¹
 - typically maintain two inner products and one matrix vector product per iteration

¹Saad 1985; Meurant 1987; Saad 1989; Chronopoulos and Gear 1989; Ghysels and Vanroose 2014.

Hiding communication

- CG is particularly sensitive to any rounding errors
- changing the order in which computations are performed can have an impact
- some of the communication hiding variants have very different behavior

Hiding communication

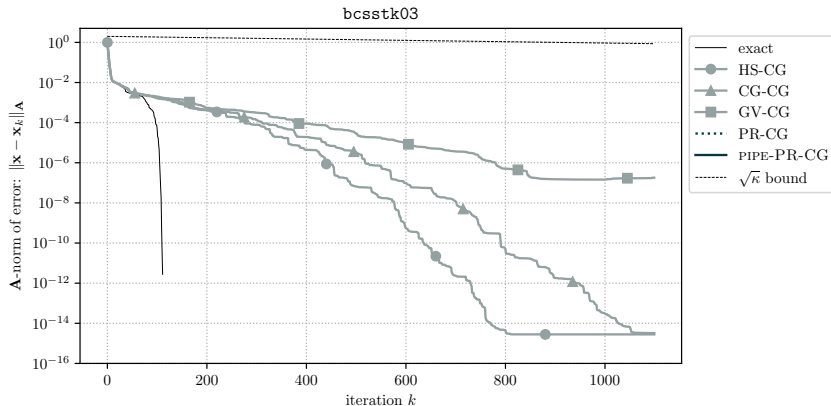
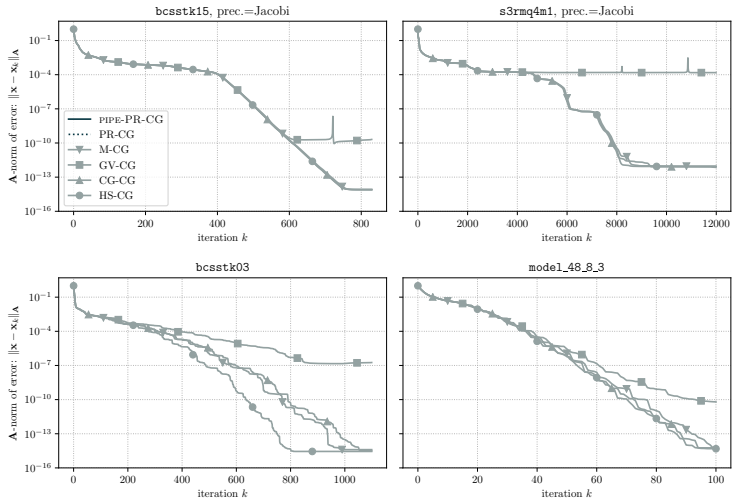


Figure: Convergence of finite precision conjugate gradient

Hiding communication



Finite precision conjugate gradient

- In finite precision orthogonality is lost, so induction based arguments for optimality of iterates no longer hold
- The **primary effects** are:
 - Delay of convergence
 - Loss of ultimately attainable accuracy

Finite precision conjugate gradient

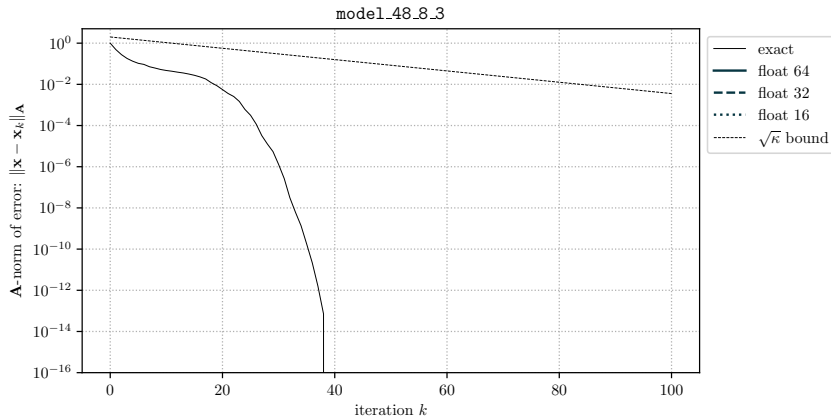


Figure: Convergence of finite precision conjugate gradient

Finite precision conjugate gradient

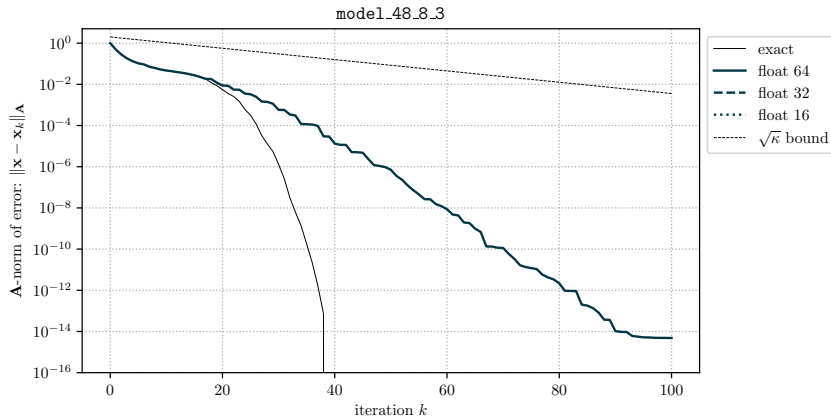


Figure: Convergence of finite precision conjugate gradient

Finite precision conjugate gradient

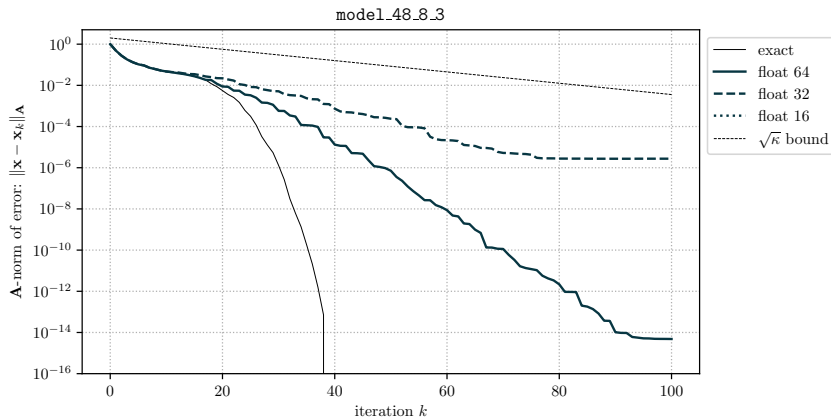


Figure: Convergence of finite precision conjugate gradient

Finite precision conjugate gradient

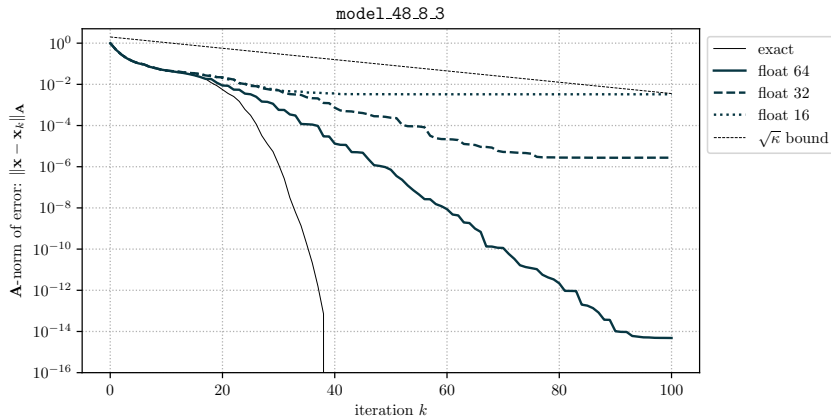


Figure: Convergence of finite precision conjugate gradient

Finite precision computations

We assume the standard model of floating point arithmetic,

$$|\text{fp}(\alpha \circ \beta) - \alpha \circ \beta| \leq \epsilon |\alpha \circ \beta| \quad (1)$$

for floating point numbers α and β and standard operations $\circ \in \{+, -, \times, \div\}$, where ϵ is the unit roundoff of the machine. From this, to first order,

$$\|\text{fp}(\mathbf{x} + \alpha \mathbf{y}) - (\mathbf{x} + \alpha \mathbf{y})\| \leq \epsilon (\|\mathbf{x}\| + 2|\alpha|\|\mathbf{y}\|) \quad (2)$$

$$\|\text{fp}(\mathbf{A}\mathbf{x}) - \mathbf{A}\mathbf{x}\| \leq \epsilon c \|\mathbf{A}\|\|\mathbf{x}\| \quad (3)$$

$$\|\text{fp}(\langle \mathbf{x}, \mathbf{y} \rangle) - \langle \mathbf{x}, \mathbf{y} \rangle\| \leq \epsilon n \|\mathbf{x}\|\|\mathbf{y}\| \quad (4)$$

where c is a dimensional constant depending on the method of matrix multiplication and n is the length of the vectors \mathbf{x} and \mathbf{y} .

Loss of Accuracy

- Most analysis uses **residual gap**: $\Delta_{\mathbf{r}_k} := (\mathbf{b} - \mathbf{A}\mathbf{x}_k) - \mathbf{r}_k$
 - if $\mathbf{r}_k \rightarrow 0$, then $\|\Delta_{\mathbf{r}_k}\| \rightarrow \|\mathbf{b} - \mathbf{A}\mathbf{x}_k\|$.
- Expression for HS-CG involves simple sum of local errors²

$$\Delta_{\mathbf{r}_k} = \Delta_{\mathbf{r}_0} + \sum_{j=1}^k (\delta_{\mathbf{r}_j} + \mathbf{A}\delta_{\mathbf{x}_k})$$

where $\delta_{\mathbf{r}_j}$ and $\delta_{\mathbf{x}_j}$ are roundoff terms from computing \mathbf{r}_j and \mathbf{x}_j in finite precision.

- similar analyses for CG-CG, GV-CG and other pipelined conjugate gradient variants³
 - CG-CG similar to HS-CG but GV-CG **amplifies** rounding errors

²Greenbaum 1989; Greenbaum 1997.

³Cools et al. 2018; Carson et al. 2018.

Perturbed Lanczos (Greenbaum 1989)

- In finite precision we can still define $\mathbf{q}_{j+1} = (-1)^j \frac{\mathbf{r}_j}{\|\mathbf{r}_j\|}$. Then we will have perturbed three term recurrence,

$$\mathbf{A}\mathbf{Q}_k = \mathbf{Q}_k\mathbf{T}_k + \gamma_k\mathbf{q}_{k+1}\mathbf{e}_k^\top + \mathbf{F}_k.$$

- If $\langle \mathbf{q}_j, \mathbf{q}_{j-1} \rangle \approx 0$ and $\|\mathbf{f}_j\| \approx 0$, then there exists a larger matrix $\bar{\mathbf{A}}$ so that
 - exact Lanczos applied to $\bar{\mathbf{A}}$ for k steps will produce \mathbf{T}_k
 - each eigenvalue of $\bar{\mathbf{A}}$ is near to some eigenvalue of \mathbf{A}
- this provides a “backwards” type analysis for Lanczos and CG

Delay of convergence

- A “good” finite precision CG implementation converges like exact CG on a larger matrix whose eigenvalues are clustered near those of \mathbf{A} .
- in this case, the variant will satisfy a minimax bound of the form,

$$\frac{\|\mathbf{e}_k\|_{\mathbf{A}}}{\|\mathbf{e}_0\|_{\mathbf{A}}} \leq \min_{p \in \mathcal{P}_k} \left[\max_{z \in \mathcal{L}(\mathbf{A})} |p(z)| \right], \quad \mathcal{L}(\mathbf{A}) = \bigcup_{i=1}^n [\lambda_i - \delta, \lambda_i + \delta].$$

- **caveat**: it remains to be proved that any commonly used CG implementations are “good”

Delay of convergence

- analysis in Greenbaum, Liu, and Chen 2019 of three term recurrence error for different variants provides some insight into why different rates of convergence are observed on some problems but not others.
- derive expressions for \mathbf{f}_k in terms of roundoff errors
 - HS-CG and CG-CG depend only on local errors, but GV-CG depends on errors made at all previous steps!
- if the previous minimax bound is particularly sensitive to the interval size (e.g. **bcsstk03**) then you see different rates of convergence

Predict-and-recompute variants

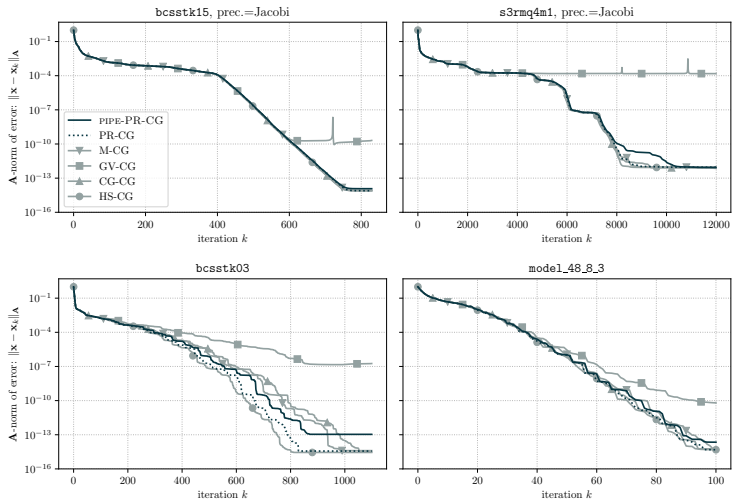
- **idea:** use recursively updated quantities as a **predictor** for their true values to allow iteration to continue, then **recompute** them directly at a later point in the iteration
- if this is done in a smart way, it won't affect the communication structure of the algorithm

Predict-and-recompute variants

variant	mem.	vec.	scal.	time
HS-CG	4 (+1)	3 (+0)	2	$2C_{\text{gr}} + T_{\text{mv}} + C_{\text{mv}}$
CG-CG	5 (+1)	4 (+0)	2	$C_{\text{gr}} + T_{\text{mv}} + C_{\text{mv}}$
M-CG	4 (+2)	3 (+1)	3	$C_{\text{gr}} + T_{\text{mv}} + C_{\text{mv}}$
PR-CG	4 (+2)	3 (+1)	4	$C_{\text{gr}} + T_{\text{mv}} + C_{\text{mv}}$
GV-CG	7 (+3)	6 (+2)	2	$\max(C_{\text{gr}}, T_{\text{mv}} + C_{\text{mv}})$
pipe-PR-M-CG	6 (+4)	5 (+3)	3	$\max(C_{\text{gr}}, T_{2\text{mv}} + C_{\text{mv}})$
pipe-PR-CG	6 (+4)	5 (+3)	4	$\max(C_{\text{gr}}, T_{2\text{mv}} + C_{\text{mv}})$

Table: Summary of costs for various conjugate gradient variants. Values in parenthesis are the additional costs for the preconditioned variants.

Predict-and-recompute variants



Predict-and-recompute variants

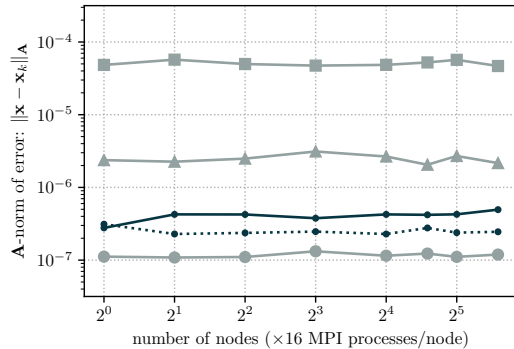
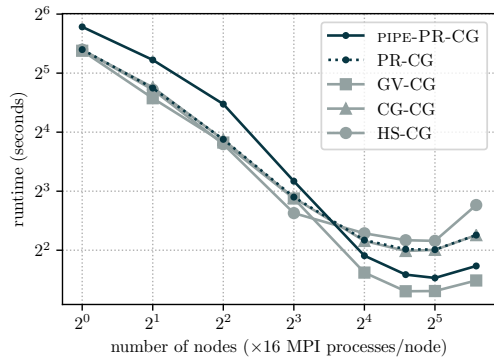


Figure: Convergence of conjugate gradient variants

Predict-and-recompute variants

- analysis in Chen and Carson 2020 of the residual gap and three term Lanczos recurrence for PR-CG and pipe-PR-CG provides insight into improved convergence
- practical use remains to be determined
- Included in PETSc v3.13: `-ksp_type pipeprcg`

Predict-and-recompute variants

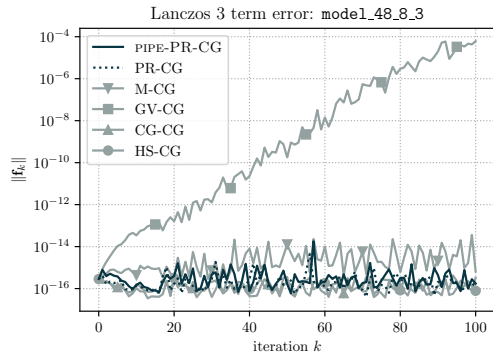
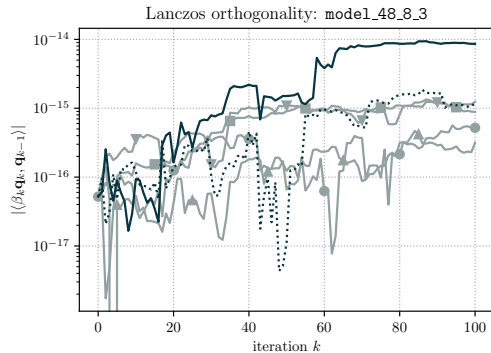


Figure: Perturbed Lanczos recurrence measures

Future work

- Try to incorporate predict-and-recompute idea into s -step methods
- Selective re-orthogonalization in low precision or high performance contexts
- Further numerical analysis of CG
 - when does $\mathbf{r}_k \rightarrow 0$?
 - when is $\langle \mathbf{r}_k, \mathbf{r}_{k-1} \rangle \approx 0$?
 - can we determine which problems will be “hard” ahead of time?

Other projects

- What effect do rounding errors have on algorithms involving randomness?
- Can finite precision be viewed as an exact computation in some space with an algebraic structure?
- Can we develop algorithms meant to perform well in finite precision without starting with an algorithm which is meant to work well in exact precision?

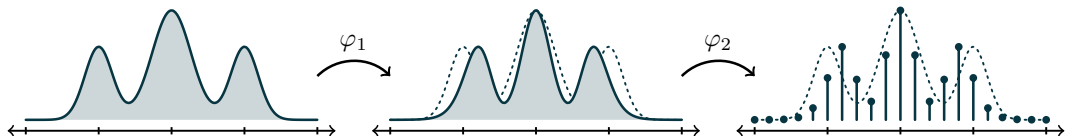


Figure: Discretizing random variable

References

- Carson, Erin C. et al. (2018). "The Numerical Stability Analysis of Pipelined Conjugate Gradient Methods: Historical Context and Methodology". In: *SIAM Journal on Scientific Computing* 40.5, A3549–A3580.
- Chen, Tyler and Erin C. Carson (2020). "Predict-and-recompute conjugate gradient variants". In: Chronopoulos, A.T. and Charles William Gear (1989). " s -step iterative methods for symmetric linear systems". In: *Journal of Computational and Applied Mathematics* 25.2, pp. 153–168.
- Cools, Siegfried et al. (Mar. 2018). "Analyzing the Effect of Local Rounding Error Propagation on the Maximal Attainable Accuracy of the Pipelined Conjugate Gradient Method". In: *SIAM Journal on Matrix Analysis and Applications* 39, pp. 426–450.
- Ghysels, Pieter and Wim Vanroose (2014). "Hiding global synchronization latency in the preconditioned Conjugate Gradient algorithm". In: *Parallel Computing* 40.7, pp. 224–238.
- Greenbaum, Anne (1989). "Behavior of slightly perturbed Lanczos and conjugate-gradient recurrences". In: *Linear Algebra and its Applications* 113, pp. 7–63.
- (1997). "Estimating the Attainable Accuracy of Recursively Computed Residual Methods". In: *SIAM Journal on Matrix Analysis and Applications* 18.3, pp. 535–551.
- Greenbaum, Anne, Hexuan Liu, and Tyler Chen (2019). *On the Convergence Rate of Variants of the Conjugate Gradient Algorithm in Finite Precision Arithmetic*.
- Meurant, Gérard (1987). "Multitasking the conjugate gradient method on the CRAY X-MP/48". In: *Parallel Computing* 5.3, pp. 267–280.
- Saad, Yousef (1985). "Practical Use of Polynomial Preconditionings for the Conjugate Gradient Method". In: *SIAM Journal on Scientific and Statistical Computing* 6.4, pp. 865–881.
- (1989). "Krylov Subspace Methods on Supercomputers". In: *SIAM Journal on Scientific and Statistical Computing* 10.6, pp. 1200–1232.

slides



`chen.pw/d/f5b`

code



`chen.pw/d/i7g`

proposal



`chen.pw/d/2mn`